

Master in Artificial Intelligence

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Advanced Human Language Technologies



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona

FIB

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Task 2.2 - DDI using neural networks

Assignment

Write a python program that parses given XML file and recognizes and classifies sentences stating drug-drug interactions. The program must use a neural network approach.

```
$ python3 ./nn-DDI.py devel.xml result.out
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e1|effect
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e2|effect
DDI-DrugBank.d211.s2|DDI-DrugBank.d211.s2.e0|DDI-DrugBank.d211.s2.e5|mechanism
...
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

General Structure

The general structure is basically the same than for the traditional ML approach:

- Two programs: one learner and one classifier.
- The learner loads the training (Train) and validation (Devel) data, formats/encodes it appropriately, and feeds it to the model, together with the ground truth.
- The classifier loads the test data, formats/encodes it in the same way that was used in training, and feeds it to the model to get a prediction.

In the case of NN, we don't need to extract features (though we **do need** some encoding)

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Input Encoding

- The input/output layers of a NN are vectors of neurons, each set to 0/1.
- Modern deep learning libraries handle this in the form of *indexes* (i.e. just provided the *position* of active neurons, omitting zeros).
- For instance, in a LSTM, each input word in the sequence may be encoded as the concatenation of different vectors each containing information about some aspect of the word (form, lemma, PoS, suffix...)
- Each vector will have only one active neuron (*one-hot encoding*), indicated by its *index*. This input is usually fed to an embedding layer.
- Our learned will need to create and store *index* dictionaries to be able to interpret the model later. See class *Codemaps* below.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Outline

Neural
Networks DDI

General
Structure

Detailed
Structure

Learner

Core task

Goals &
Deliverables

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Learner - Main program

Neural
Networks DDI
General
Structure
Detailed
Structure
Learner
Core task
Goals &
Deliverables

```
1 def train(trainfile, validationfile, params, modelname) :
2     ## learns a NN model using trainfile as training data, and validationfile
3     ## as validation data. Saves learnt model in a file named modelname
4     # load pickle datasets (or parse if needed)
5     traindata = Dataset(trainfile)
6     valdata = Dataset(valfile)
7
8     # create indexes from training data
9     codes = Codemaps(traindata, params)
10    # encode datasets
11    train_loader = encode_dataset(traindata, codes, params)
12    val_loader = encode_dataset(valdata, codes, params)
13
14    # build network
15    network = ddiCNN(codes)
16
17    # save indexs
18    os.makedirs(modelname, exist_ok=True)
19    torch.save(network, os.path.join(modelname, "network.nn"))
20    codes.save(os.path.join(modelname, "codemaps"))
21    # train each epoch, keep the best model on validation
22    best = 0
23    for epoch in range(params["epochs"]):
24        train(network, epoch, train_loader)
25        acc = validation(network, val_loader)
26        if acc>best :
27            best = acc
28            torch.save(network, os.path.join(modelname,f"network.nn"))
```

Outline

Neural
Networks DDI

General
Structure

Detailed
Structure

Classifier

Core task

Goals &
Deliverables

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Classifier - Main program

```
1 def predict(modelname, datafile, params, outfile) :
2     """
3         Loads a NN model from file modelname and uses it to extract
4         drug interactions form datafile
5     """
6     # Load model
7     model = torch.load(os.path.join(modelname, "network.nn"),
8                         map_location=torch.device(used_device))
9     model.eval()
10    # load indexes
11    codes = Codemaps(os.path.join(modelname, "codemaps"), params)
12    # load data to classify
13    testdata = Dataset(datafile)
14    test_loader = encode_dataset(testdata, codes, params)
15
16    # run each example and obtain prediction
17    Y = []
18    for X in test_loader:
19        # X is a list of input tensors (no labels were loaded in the
20        # dataloader)
21        y = model.forward(*X) # run example through the network
22        # add results to result list
23        Y.extend([codes.idx2label(torch.argmax(s)) for s in y])
24
25    # output results
26    output_interactions(testdata, Y, outfile)
```

NOTE: Observe the output structure (one class per sentence+pair), different from the NER task (one class per token).

Neural
Networks DDI

General
Structure

Detailed
Structure

Classifier

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - parse_data

Processing the whole dataset with Spacy takes some time, so it is convenient to run it once and for all:

```
$ python3 run.py parse
```

This will create pickle files for train and devel data in preprocessed folder, so they can be used by the feature extractor.

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - Dataset

```
1 class Dataset:  
2     ## constructor:  
3     ## If 'filename' is a '.pck' file, load data set pickle file  
4     ## Otherwise, assume 'filename' is an xml file: Tokenize  
5     ## each sentence, and store a list of (sentence, entity pair).  
6     ## For each (sentence, entity pair), store token information,  
7     ## masking all entities in the sentence.  
8     def __init__(self, filename)  
9  
10    ## saves dataset to a piclke file (to avoid repeating parsing)  
11    def save(self, filename)  
12  
13    ## iterator to get sentences in the dataset  
14    def sentence(self)  
15    ...
```

Class Dataset will *mask* the target entities in the input sentence:

Original sentence: *Exposure to oral ketamine is unaffected by itraconazole compounds but greatly increased by ticlopidine.*

Pair	Masked sentence
e0-e1	Exposure to oral DRUG1 is unaffected by DRUG2 but greatly increased by DRUG_OTHER.
e0-e2	Exposure to oral DRUG1 is unaffected by DRUG_OTHER but greatly increased by DRUG2.
e1-e2	Exposure to oral DRUG_OTHER is unaffected by DRUG1 but greatly increased by DRUG2.

Auxiliary classes - Codemaps

```
1 class Codemaps :
2     # Constructor: create code mapper either from training data, or
3     #           loading codemaps from given file.
4     #           If 'data' is a Dataset, and lengths are not None,
5     #           create maps from given data.
6     #           If data is a string (file name), load maps from file.
7     def __init__(self, data, maxlen=None, suflen=None)
8     # Save created codemaps in file named 'name'
9     def save(self, name)
10    # Convert a Dataset into lists of word codes and sufix codes
11    # Adds padding and unknown word codes.
12    def encode_words(self, data)
13    # Convert the gold labels in given Dataset into a list of label codes.
14    # Adds padding
15    def encode_labels(self, data)
16    # get word index size
17    def get_n_words(self)
18    # get suf index size
19    def get_n_sufs(self)
20    # get label index size
21    def get_n_labels(self)
22    # get index for given word
23    def word2idx(self, w)
24    # get index for given suffix
25    def suff2idx(self, s)
26    # get index for given label
27    def label2idx(self, l)
28    # get label name for given index
29    def idx2label(self, i)
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

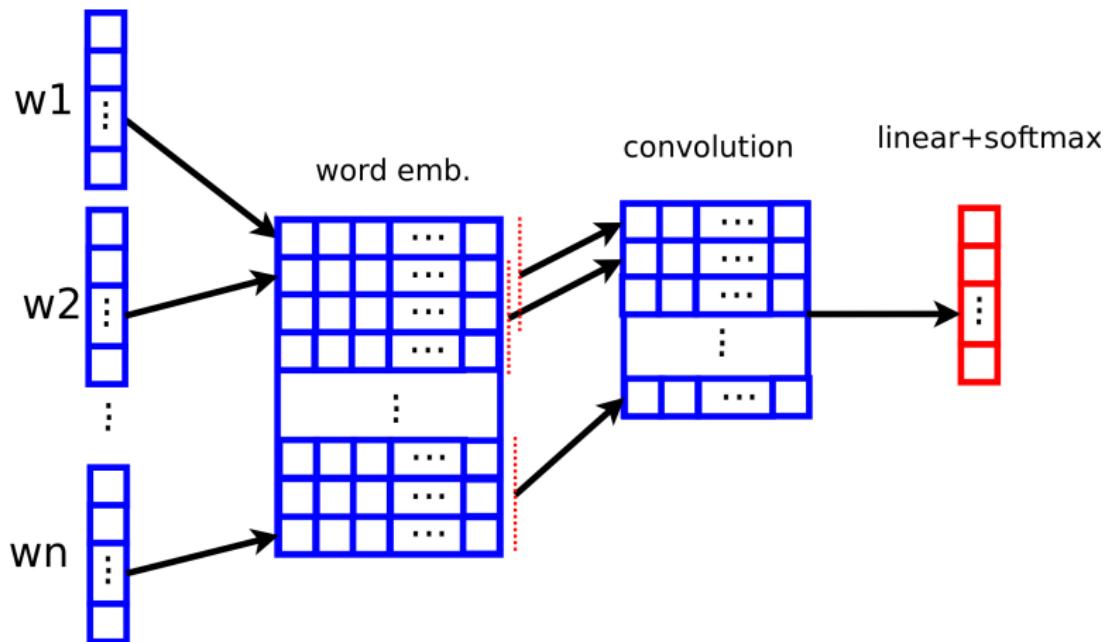
Required functions - network.py

Neural
Networks DDI
General
Structure
Detailed
Structure
Auxiliary classes
Core task
Goals &
Deliverables

```
1 class ddiCNN(nn.Module):
2
3     def __init__(self, codes) :
4         super(ddiCNN, self).__init__()
5         # get sizes from index
6         n_words = codes.get_n_words()
7         n_labels = codes.get_n_labels()
8         max_len = codes.maxlen
9         # create embedding layer
10        embW_sz = 100
11        self.embW = nn.Embedding(n_words, embW_sz, padding_idx=0)
12        # create CNN layer
13        cnn_out_sz = 32
14        self.cnn = nn.Conv1d(embW_sz, cnn_out_sz, kernel_size=2, stride=1,
15                           padding='same')
16        self.drop2 = nn.Dropout(0.2)
17        # final classification layer
18        self.out = nn.Linear(cnn_out_sz*self.max_len, n_labels)
19
20    def forward(self, w):
21        # run layers on given data
22        x = self.embW(w)          # apply embedding layer
23        x = x.permute(0,2,1)      # set shape appropriate for CNN input
24        x = self.cnn(x)           # apply CNN
25        x = func.relu(x)          # activation function
26        x = x.flatten(start_dim=1) # set shape appropriate for linear input
27        x = self.out(x)           # final classification layer
28
29    return x
```

Network architecture

Neural Networks DDI
General Structure
Detailed Structure
Auxiliary classes
Core task
Goals & Deliverables



Outline

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Build a good NN-based DDI detector

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

- DDI is **not** a sequence tagging task (which assign one label per word), but a sentence classification, where a single label is assigned to the whole sentence (or sentence + entity pair in this case).
- Good results may be achieved using a CNN, as in the provided example.
- The problem also may be approached with an LSTM. Note that instead of getting the output at each word, only the output at the end of the sequence must be used (or the output of all words must be combined to feed further layers).
- It is also possible to combine LSTM and CNN layers.
- If you add extra input layers (e.g. lemma, pos, lowercase word, etc) you will need to add one embedding layer after the input, that is where the created indexes will become handy.
- You may get inspiration for an architecture from these examples: [1], [2], [3], [4], some of the papers provided in labAHLT package in papers/SharedTask/otherSystems, or just googling for *semeval DDI neural networks*.

Build a good NN-based DDI detector

Strategy: Experiment with different NN architectures and possibilities.

Some elements you can play with:

- Embedding dimensions, number and kind of layers, used optimizer...
- Using just CNN, just a LSTM, or a LSTM+CNN combination
- Using lowercased and/or non lowercased word embeddings
- Initialitzing embeddings with available pretrained model
- Using extra input (e.g. lemma embeddings, PoS embeddings, suffix/prefix embbedings, ...)
- Adding extra dense layers, with different activation functions
- Using pretrained transformers such as Bert as the first layers of your network.
- Adding attention layers
- ...etc.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Build a good NN-based DDI detector

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Warnings:

- Neural Network training uses randomization, so different runs of the same program will produce different results. For repeatable results, use a random seed (and/or run the training several times).
- During training, *accuracy* on training and validation sets is reported. Those values are usually over 85%. However, this is due to the fact that most of the pairs have label “null” (no interaction). Accuracy values around 85% correspond to very low F_1 values. To get a reasonable F_1 , validation set accuracy should reach about 89-90%.

To precisely evaluate how your model is doing, **do not rely** on reported accuracy: run the classifier on the development set and use the evaluator.

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Exercise Goals

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

What you should do:

- Experiment with architecture variations.
- Experiment with different learning hyperparameters.
- Experiment with different input information
- Keep track of tried variants and parameter combinations.

What you should **NOT** do:

- Get insights about errors from `devel` dataset. You have `train` for that. `devel` is only used to evaluate the performance of a given configuration.
- Select architectures or hyperparameters based on system performance on `test` dataset. You have `devel` for that. `test` is only used to evaluate model generalization ability once the best configuration has been chosen.

Deliverables

At the end of the DDI task, you will need to deliver a single report on the work carried out on DDI-ML, DDI-NN, and DDI-LLM systems

So, during the development and experimentation on DDI-NN:

- keep track of tried/discard architectures/hyperparameters
- keep track of tried/discard input information.
- Record obtained results in the different experiments, and compile the information you'll later need to elaborate the report.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables