

# Master in Artificial Intelligence

## Advanced Human Language Technologies

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables



UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

---

Facultat d'Informàtica de Barcelona

FIB

# Outline

## 1 LLM Drug-drug Interaction

### 2 General Structure

### 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

### 4 Core task

### 5 Goals & Deliverables

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Task 2.3 - DDI using large language models

## Assignment

Write a python program that parses the given XML file and recognizes and classifies sentences stating drug-drug interactions. The program must use a LLM-based approach.

```
$ python3 ./nn-DDI.py devel.xml result.out  
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e1|effect  
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e2|effect  
DDI-DrugBank.d211.s2|DDI-DrugBank.d211.s2.e0|DDI-DrugBank.d211.s2.e5|mechanism  
...
```

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

### 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# General Structure

We will use the same representation than in the NN approach:

Identify the target pair of drugs in a sentence with labels **DRUG1** and **DRUG2**, and ask the model to decide whether the sentence states an interaction between them.

Original sentence:

*Exposure to oral **ketamine** is unaffected by **itraconazole compounds** but greatly increased by **ticlopidine**.*

Generated examples (potential interactions)

Pair	Masked sentence
e0-e1	Exposure to oral <b>DRUG1</b> is unaffected by <b>DRUG2</b> but greatly increased by <b>DRUG_OTHER</b> .
e0-e2	Exposure to oral <b>DRUG1</b> is unaffected by <b>DRUG_OTHER</b> but greatly increased by <b>DRUG2</b> .
e1-e2	Exposure to oral <b>DRUG_OTHER</b> is unaffected by <b>DRUG1</b> but greatly increased by <b>DRUG2</b> .

# General Structure

- We will need to convert the dataset to/from the pseudo-XML representation.
- We will need to instruct the LLM with a suitable prompt describing the task, followed by each target sentence to annotate.
- We will be using two different strategies:
  - **Few-shot learning:** Add some solved examples in the prompt, so the model understands better what is expected.
  - **Fine-tuning:** Tune the LLM weights using several (dozens, hundreds, ...) solved examples, so it learns the task and does not need the few-show examples each time.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Few-shot learning

- The LLM is asked to perform an specific task in a prompt, which contains:
  - The task description. It must be precise and well explained.
  - Some pairs of (input, expected output) examples.
  - The input for the example to solve
- Adjustable elements:
  - The **prompt**: How the target task is described. Includes specific orders to produce output in certain ways/formats, and NOT to produce undesired output.
  - The **number** of few-shot examples.
  - The **variety** of used few-shot examples (e.g. making sure all interaction types appear in some example).

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Fine-Tuning

- The LLM is presented a collection of prompts, and the model weights are iteratively updated. Each prompt includes:
  - The task description.
  - One pair of (input, expected output) example.
- Adjustable elements:
  - The **prompt**: How required task is described. Specific orders to produce output in certain ways, and NOT to produce undesired output. Not as crucial as in few-shot
  - The **amount** of fine tuning examples
  - The **distribution** of used examples (e.g. making sure all interaction types are represented enough, maybe balancing too frequent or too scarce classes).

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

## 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Few shot learning

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

## 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# Few shot - Annotator

```
1 # load training data (FS examples)
2 trainfile = os.path.join(paths.DATA,traindata+".xml")
3 fs_examples = Examples(trainfile, "DDI").select_examples(num_few_shot,
4                                                               balanced=True)
5 # load prompts, create few-shot prompt
6 prompts = Prompts(promptfile, fs_examples)
7
8 # load test data
9 testfile = os.path.join(paths.DATA,testdata+".xml")
10 test = Examples(testfile, "DDI")
11
12 # load model and tokenizer
13 if ollama:
14     engine = Inference(model, ollama=True)
15 else :
16     MODEL_PATH = f"/scratch/nas/1/PDI/mml0/models/{model}"
17     engine = Inference(MODEL_PATH, quantized=quantized)
18
19 # annotate each example intestdata
20 annotated = []
21 for i,ex in enumerate(test.select_examples()):
22     # create prompt for this example, adding it to FS prompt
23     messages = prompts.prepare_messages(ex['input'])
24     # call model to generate response
25     gen_text = engine.generate(messages)
26     # store responses
27     ex['predicted'] = gen_text
28     ex['evaluator'] = test.eval_format(ex,gen_text)
29     annotated.append(ex)
30
31 # print results
32 for e in annotated: print(e["evaluator"])
```

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Few shot learning

Core task

Goals &  
Deliverables

# Outline

## 1 LLM Drug-drug Interaction

### 2 General Structure

### 3 Detailed Structure

- Few shot learning
- **Fine-tuning**
- Auxiliary classes

### 4 Core task

### 5 Goals & Deliverables

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Fine-tuning

Core task

Goals &  
Deliverables

# Fine tuning - Trainer

```
1 # load prompts
2 prompts = Prompts(promptfile)
3
4 # load model and tokenizer
5 MODEL_PATH = f"/scratch/nas/1/PDI/mml0/models/{model}"
6 engine = FineTuning(MODEL_PATH, quantized=quantized)
7
8 # load and tokenize datasets
9 trainfile = os.path.join(paths.DATA,traindata+".xml")
10 train_examples = Examples(trainfile, "DDI").select_examples(5000,
11                                         balanced=True)
12 train_dataset = engine.tokenize_dataset(train_examples, prompts)
13
14 valfile = os.path.join(paths.DATA,validata+".xml")
15 val_examples = Examples(valfile, "").select_examples(500,
16                                         balanced=True)
17 val_dataset = engine.tokenize_dataset(val_examples, prompts)
18
19 # Fine-tune the model and save results
20 engine.train(train_dataset,
21             val_dataset,
22             outputdir)
```

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Fine-tuning

Core task

Goals &  
Deliverables

# Fine tuning - Annotator

LLM  
Drug-drug  
Interaction  
  
General  
Structure  
  
Detailed  
Structure  
  
Fine-tuning  
  
Core task  
  
Goals &  
Deliverables

```
1 # load prompts
2 prompts = Prompts(promptfile)
3
4 # load test/devel dataset
5 testfile = os.path.join(paths.DATA,testdata+".xml")
6 test = Examples(testfile, "DDI")
7
8 # load model and tokenizer
9 MODEL_PATH = f"/scratch/nas/1/PDI/mml0/models/{model}"
10 engine = Inference(MODEL_PATH, quantized=quantized, peft=weightdir)
11
12 # analyze each example
13 annotated = []
14 for i,ex in enumerate(test.select_examples()):
15     # prepare sequence of messages for this example
16     messages = prompts.prepare_messages(ex['input'])
17     # call model to generate response
18     gen_text = engine.generate(messages)
19     # extract json from response
20     ex["predicted"] = gen_text
21     ex['evaluator'] = test.eval_format(ex,gen_text)
22     annotated.append(ex)
23
24 # save output
25 for e in annotated: print(e["evaluator"])
```

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Auxiliary classes

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

## 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# Auxiliary classes - Examples

```
1 class Examples:  
2     ## constructor: parses data XML file, and converts sentences into  
3     ## the format expected by the LLM for the given task (NER/DDI).  
4     def __init__(self, xmlfile, task)  
5  
6     ## get 'numFS' examples (or all dataset if omitted)  
7     ## balance classes if requested  
8     def select_examples(self, numFS=-1, balanced=False)  
9  
10    ## convert LLM output into the format expected by the  
11    ## evaluator for the task this instance is loaded for  
12    def eval_format(self, ex, text)
```

- Class `Examples` takes care of loading the dataset XML files and converting the examples into the appropriate representation for the LLM.
- Class `Examples` also converts LLM responses into the output format expected by the evaluator.
- For DDI, `balanced` parameter helps selecting training examples.
- Method `select_examples` can be replaced with a custom selection strategy.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Auxiliary classes

Core task

Goals &  
Deliverables

# Auxiliary classes - Prompts

```
1 class Prompts:  
2     # Constructor: loads prompts from json file (sysprompt and usrprompt)  
3     # If a list of fs_examples is provided, they are added to the prompt  
4     def __init__(self, promptfile, fs_examples=[])  
5  
6         # Build prompt for a particular example, concatenating the  
7         # base prompt (sysprompt + usrprompt + fs_xamples --if any)  
8         # with the given user request 'question', plus the expected  
9         # answer --if provided  
10        def prepare_messages(self, question, answer="")
```

- The constructor stores a prompt that contains the system prompt and the user prompt. For few-shot, it also contains the provided solved examples.
- Method `prepare_messages` is used in few-shot to add a request (question) to the prompt (which already has some solved examples), without the corresponding answer.
- Method `prepare_messages` is used in fine-tuning training to add a complete solved example (request plus expected answer). In this scenario, the base prompt does not contain few-shot examples.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Auxiliary classes

Core task

Goals &  
Deliverables

# Auxiliary classes - model.Inference

```
1 class Inference:  
2     # load given model in inference mode  
3     def __init__(self, model_path, quantized=False, peft=None, ollama=False)  
4  
5     # Generate completion for given messages  
6     def generate(self, messages)  
7
```

- Constructor loads requested model (either using ollama or HF transformers) in inference mode. The model can be loaded in full, or quantized to 4 bits. If `peft` is given it is expected to be the path to fine-tuned LoRa adapters for the given model, which are also loaded.
- Method `generate` uses the loaded LLM to generate a completion for given prompt (typically, the answer to the last request in it).

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Auxiliary classes

Core task

Goals &  
Deliverables

# Auxiliary classes - model.FineTuning

```
1 class Inference:  
2     # load given model in LoRa fine tuning mode.  
3     def __init__(self, model_path, quantized=False)  
4  
5         # Tokenizes given dataset with the model appropriate tokenizer.  
6         # Returns a "datasets.Dataset" object suitable for HF trainer  
7         def tokenize_dataset(self, dataset, prompts)  
8  
9             # Fine tunes model with given train and validation data.  
10            # Saves tuned weights in outputdir  
11            def train(self, train_dataset, val_dataset, outputdir)  
12
```

- The constructor loads requested model (quantized if required) and adds LoRa adapters to be tuned.
- Method `tokenize_datasets` uses model tokenizer to preprocess the training/validation data and convert them into batched `datasets.Dataset`, suitable for the trainer.
- Method `train` performs the fine-tuning using training data and computes loss on validation data at each epoch.
- Method `train` can be adjusted to change learning parameters if needed.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Auxiliary classes

Core task

Goals &  
Deliverables

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

## 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# Build a good LLM-based DDI detection

**Strategy:** Experiment with different **few-shot** configurations:

- Prompts:
  - Explain better/more clearly the task.
  - If the model produces wrong output, insist on what is the expected output or its right format.
  - If the model misses some expected output, insist on its importance.
  - If the model produces unrequested output, explicitly request it not to do so.
- Few-shot examples
  - Experiment with different number of few-shot examples.
  - Experiment with better example selection methods (even hand-picking most informative examples from training set). Modify `select_examples` method to achieve this.
- Model
  - Use different models (llama, qwen, gemma, mistral... ).
  - Experiment with/without quantization

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Build a good LLM-based DDI detection

**Strategy:** Experiment with different **fine-tuning** configurations:

- Prompts:
  - Improve prompts as in few-shot (although it has much less effect here).
- Training examples
  - Experiment with different number of training examples.
  - Experiment with better example selection methods.  
Modify `select_examples` method to achieve this.
- Model
  - Use different models (llama, qwen, gemma, mistral... ) and sizes.
  - Experiment with/without quantization
- Learning hyperparameters
  - Learning rate
  - Epochs
  - batch size (or gradient accumulation steps)

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Build a good LLM-based DDI detection

## Few-shot caveats:

- **Model context size:** Too many few-shot examples might fill the model context window and make it forget the instructions.
- **Available GPU RAM:** Too many few-shot examples, even if they don't fill the context window, may not fit into GPU RAM.
- **Token offset errors:** The model might slightly alter the text, which will result in wrong offsets when converting the XML tags to the format expected by the evaluator.

## Fine-tuning caveats:

- **Available GPU RAM:** Fine tuning a model takes about 3 times the RAM used in inference. Bigger GPUs are required. Smaller batch sizes may help. Quantization too.
- **Available time:** Fine tuning takes a lot of time, since the dataset needs to be run through the model, loses computed, weights updated. For each example, several epochs.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Outline

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## 1 LLM Drug-drug Interaction

## 2 General Structure

## 3 Detailed Structure

- Few shot learning
- Fine-tuning
- Auxiliary classes

## 4 Core task

## 5 Goals & Deliverables

# Exercise Goals

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

What you should do:

- Experiment with few-shot and fine-tuning variations.
- Experiment with different prompts and few-shot examples.
- Experiment with different prompts and training examples.
- Experiment with different learning hyperparameters.
- Keep track of tried variants and parameter combinations.

What you should **NOT** do:

- Get insights about errors from `devel` dataset. You have `train` for that. `devel` is only used to evaluate the performance of a given configuration.
- Select architectures or hyperparameters based on system performance on `test` dataset. You have `devel` for that. `test` is only used to evaluate model generalization ability once the best configuration has been chosen.

# Deliverables

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

At the end of the DDI task, you will need to deliver a single report on the work carried out on DDI-ML, DDI-NN, and DDI-LLM systems

So, during the development and experimentation on DDI-LLM:

- keep track of tried/discard architectures/hyperparameters
- keep track of tried/discard input information.
- Record obtained results in the different experiments, and compile the information you'll later need to elaborate the report.

## Annex: Resources to use LLMs. Ollama

ollama can be used for few-shot experiments.

- Ollama runs LLMs with or without GPU (much slower without)
- Install ollama: <https://ollama.com/download>
- Install ollama python API: pip install ollama
- Download desired ollama models
  - ollama pull llama3.2:3b
  - ollama pull qwen3:8b
  - ollama pull ministral3:8b
- See <https://ollama.com/library> for a list
- Use provided `fewshot.py` with `-ollama` option.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## Annex: Resources to use LLMs. Huggingface

Huggingface Transformers can be used for fine-tuning experiments.

- If you have a GPU (with at least 8-10GB), you can install requirements.txt in a python virtual environment, and run the programs out-of-the box.
- Use HF model id (e.g. `meta-llama/Llama-3.2-3B-Instruct`) as MODEL\_PATH when loading the model.
- If you don't have a GPU, you can use [google colab](#), or UPC [boada](#) cluster, where llama3.2-3B and qwen3-3B are already installed.

LLM  
Drug-drug  
Interaction

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables